

# Artificial Intelligence

## **Lecture 6:**

## **Introduction to Reinforcement Learning**


Henrique Lopes Cardoso, Luís Paulo Reis

hlc@fe.up.pt, lpreis@fe.up.pt



# What is Reinforcement Learning?

- **Reinforcement Learning (RL)** is focused on goal-directed learning from interaction
- RL is **learning what to do** – how to map situations to actions – so as to maximize a numerical **reward** signal
  - The learner is not told which actions to take: it must discover which actions yield the most reward by trying them
  - Typically, actions may affect not only immediate reward but also the next situation and subsequent rewards
- The **exploration-exploitation** tradeoff
  - Agent must prefer actions that it knows to be effective – *exploit*
  - But to discover such actions, it has to try actions not selected before – *explore*



# RL vs (Un)Supervised Learning

- Different from **supervised learning**
  - In interactive problems it is impractical to obtain examples of desired behavior
  - In uncharted territory, an agent must learn from its own experience
- Different from **unsupervised learning**
  - RL is trying to maximize a reward signal, not trying to find hidden structure in collections of unlabeled data
- RL explicitly considers the *whole* problem of a **goal-directed agent interacting with an uncertain environment**
  - Creating a behavior model while applying it in the environment
- RL is the closest form of ML to the kind of learning humans do


# Learning to Play Tic-Tac-Toe

- Rule-based approach
  - Need to hardcode rules for each possible situations that might arise in a game
- Minimax
  - Assumes a particular way of playing by the opponent
- Dynamic programming can compute an optimal solution for any opponent
  - But requires as input a complete specification of that opponent (state/action probabilities)
- Can we obtain such information from **experience**?
  - Play many games against the opponent!

X	O	O
O	X	X
		X

# Learning to Play Tic-Tac-Toe

- **States**
  - Possible configurations of the board
- **Actions**
  - Possible moves to make
- **Policy**
  - Which action should I play in each state?
- **Reward**
  - How good was the chosen action?



# Elements of RL

- **Policy  $\pi$**

- How should the agent behave over time?
- A policy is a (possibly stochastic) **mapping from perceived states to actions**

- **Reward signal  $r$**

- Defines the **goal** of the RL problem
- On each time step, the environment sends a **reward** to the RL agent – the agent's goal is to **maximize the total reward** received over the long run

- **Value function  $v$**

- Specifies what is good in the long run
- The **value** of a state is the **total amount of reward** an agent can expect to accumulate from that state onwards (it takes into account future rewards)

→ We seek actions that bring about states of **highest value**, not highest reward, because these actions obtain the greatest amount of reward over the long run


# Bandit Problems

- A simple setting with a **single state**



- $K$ -armed bandit problem
  - There are  $k$  different **actions**
  - After each action a numerical reward is received from a stationary probability distribution
  - Each action has a **value** – its expected or mean reward, not known by the agent:  $q_*(a) \doteq \mathbb{E}[R_t | A_t = a]$
  - The agent **estimates**, at time step  $t$ , the value of an action  $a$ :  $Q_t(a)$

# Bandit Problems



- Selecting *greedy* actions (whose estimated value is greatest): **exploiting**
- Selecting non-greedy actions: **exploring**
  - Improve estimates of non-greedy actions' value
- Lower reward in the short run (during *exploration*), but higher in the long run – after discovering the best actions, we can *exploit* them many times



# Estimating Action Values

- Sample average:

$$Q_t(a) \doteq \frac{\sum_{i=1}^{t-1} R_i \cdot 1_{A_i=a}}{\sum_{i=1}^{t-1} 1_{A_i=a}}$$

- Update rule:

$$Q_{n+1} = Q_n + \frac{1}{n} [R_n - Q_n]$$

$$NewEstimate \leftarrow OldEstimate + StepSize [Target - OldEstimate]$$

- The target indicates a desirable direction in which to move
  - The *step-size parameter* changes from time step to time step
- 
- Giving more weight to recent rewards – *constant step-size parameter*:

$$Q_{n+1} \doteq Q_n + \alpha [R_n - Q_n]$$

where  $\alpha \in (0,1]$




# Action Selection

- *Greedy* action selection (always exploits):  $A_t \doteq \underset{a}{\operatorname{argmax}} Q_t(a)$
- *$\epsilon$ -greedy* action selection: behave greedily most of the time, but with small probability  $\epsilon$  select randomly from among all the actions
  - $Q_t(a)$  will converge to  $q_*(a)$  if  $a$  is selected sufficiently often
- *Soft-max* action selection (Boltzmann distribution):

$$\Pr\{A_t = a\} \doteq \frac{e^{Q_t(a)/\tau}}{\sum_{b=1}^k e^{Q_t(b)/\tau}}$$

where  $\tau$  is a temperature parameter: if high, actions will tend to be equiprobable; if low, action values matter more; if  $\tau \rightarrow 0$ , then we have greedy action selection


# The 10-armed Testbed



# Markov Decision Processes

- In the general setting we have **many states**
- **Markov Decision Processes (MDP)** are a classical formalization of sequential decision making
  - Actions influence not just immediate rewards, but also subsequent situations (states) and thus future rewards
- In a finite MDP, there is a finite number of states, actions and rewards
- In MDPs we estimate the value  $q_*(s, a)$

# Agent-Environment Interface



- Dynamics of the MDP:

$$p(s', r | s, a) \doteq \Pr\{S_t = s', R_t = r \mid S_{t-1} = s, A_{t-1} = a\}$$

- The probability of each possible value for  $s'$  and  $r$  depends only in the immediately preceding state  $s$  and action  $a$
- The state must include all relevant information about the past agent-environment interaction – **Markov property**


# Example: Recycling Robot

- A robot has to decide whether it should (1) actively search for a can, (2) wait for someone to bring it a can, or (3) go to home base and recharge
- Searching is better (higher probability of getting a can) but runs down battery; if out of battery, the robot has to be rescued
- Decisions made on the basis of current energy level: high, low
- Reward is zero except when getting a can, and negative if out of battery

$$\begin{aligned}\mathcal{S} &= \{high, low\} \\ \mathcal{A}(high) &= \{search, wait\} \\ \mathcal{A}(low) &= \{search, wait, recharge\} \\ r_{search} &> r_{wait}\end{aligned}$$

$s$	$a$	$s'$	$p(s'   s, a)$	$r(s, a, s')$
high	search	high	$\alpha$	$r_{search}$
high	search	low	$1 - \alpha$	$r_{search}$
low	search	high	$1 - \beta$	$-3$
low	search	low	$\beta$	$r_{search}$
high	wait	high	1	$r_{wait}$
high	wait	low	0	-
low	wait	high	0	-
low	wait	low	1	$r_{wait}$
low	recharge	high	1	0
low	recharge	low	0	-

# Example: Recycling Robot



# Goals and Rewards

- A **reward signal** is used to define the **goal** of the agent
    - Learning to walk: reward proportional to the robot's forward motion
    - Learning to Escape from a maze: reward -1 for any state prior to escape (encourage escaping as quickly as possible)
    - Learning to find empty cans for recycling: reward of 0 most of the time, +1 for each can collected
    - Learning to play checkers or chess: reward +1 for winning, -1 for losing, and 0 for drawing and nonterminal positions
  - Provide **rewards** in such a way that by **maximizing** them the agent will also achieve our **goal**
    - The agent's goal is to **maximize the cumulative reward** it receives in the long run
- The reward signal is a way of communicating to the robot **what** you want it to achieve, not **how**



# Returns and Episodes

- Agent wants to maximize **expected return**

$$G_t \doteq R_{t+1} + R_{t+2} + R_{t+3} + \dots + R_T$$


- **Episodic tasks**: when the agent-environment interaction breaks naturally into subsequences – **episodes**
  - From a **starting state** to a **terminal state**
  - Followed by a reset to another starting state, chosen independently of how the previous episode ended
- **Continuing tasks** do not break naturally into identifiable episodes (*e.g.*, on-going process-control)
  - Problem with calculating  $G_t$ :
    - $T = \infty$
    - $G_t$  could also be infinite (if rewards are positive at each time step)

# Returns and Episodes

- Adding **discounting**: agent wants to maximize **expected discounted return**


$$G_t \doteq R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \qquad G_t \doteq R_{t+1} + \gamma G_{t+1}$$

- $0 \leq \gamma \leq 1$  is the **discount rate**
  - If  $\gamma = 0$  the agent is “myopic” (only immediate reward matters)
  - As  $\gamma$  approaches 1, the agent becomes more farsighted (strongly considers future rewards)
- $G_t$  is now finite, even if summing an infinite number of terms
- Applicable also to episodic tasks, if we consider a final absorbing state:



# Example: Pole Balancing

- Move a **cart** so as to keep a **pole** from falling over
  - Failure if the pole falls past a given angle or if the cart runs off the track
  - The pole is reset to vertical after each failure



- **Episodic task**: reward +1 except when failure
  - return is the number of steps until failure
- **Continuing task**, using discounting: reward  $-1$  on each failure and 0 otherwise
  - return is  $-\gamma^K$ , where  $K$  is the number of steps before failure


# Policies and Value Functions

- RL algorithms involve estimating **value functions**
  - How good (in terms of expected return) is it to be in a given state?
  - How good is it to perform a given action in a given state?
- Future rewards depend on the choice of actions
  - Value functions are defined with respect to **policies** (ways of acting)
- **Policy**: a mapping from states to probabilities of selecting each possible action
  - $\pi(a|s) = \Pr(A_t = a|S_t = s)$

→ RL methods specify how the policy (*i.e.*, the probability distribution over  $a \in \mathcal{A}(s)$  for each  $s \in \mathcal{S}$ ) is changed with experience

# Policies and Value Functions

- **State-value** function  $v_{\pi}(s)$ 
  - Expected return when starting in  $s$  and following  $\pi$  thereafter
  - $v_{\pi}(s) \doteq \mathbb{E}_{\pi}[G_t | S_t = s]$  (*Bellman equation*)
- **Action-value** function  $q_{\pi}(s, a)$ 
  - Expected return when taking action  $a$  in state  $s$ , and following  $\pi$  thereafter
  - $q_{\pi}(s, a) \doteq \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$
- The value functions  $v_{\pi}$  and  $q_{\pi}$  can be estimated from experience
- Example: using a random policy, with  $\gamma = 0.9$ :



Gridworld



- Off-grid actions have no effect, with  $r = -1$
- Any action from A gets to A', with  $r = +10$
- Any action from B gets to B', with  $r = +5$


3.3	8.8	4.4	5.3	1.5
1.5	3.0	2.3	1.9	0.5
0.1	0.7	0.7	0.4	-0.4
-1.0	-0.4	-0.4	-0.6	-1.2
-1.9	-1.3	-1.2	-1.4	-2.0

$v_{\pi}$

# Optimal Policy and Value Function


- **Optimal policy  $\pi_*$** : expected return is greater than any other policy

- **Optimal state-value function**:  $v_*(s) \doteq \max_{\pi} v_{\pi}(s)$
- **Optimal action-value function**:  $q_*(s, a) \doteq \max_{\pi} q_{\pi}(s, a)$



- Once we know  $v_*$  or  $q_*$ , the **optimal policy is greedy**

- Example:



Gridworld



Actions


22.0	24.4	22.0	19.4	17.5
19.8	22.0	19.8	17.8	16.0
17.8	19.8	17.8	16.0	14.4
16.0	17.8	16.0	14.4	13.0
14.4	16.0	14.4	13.0	11.7

$v_*$

→	↕	←	↕	←
↙	↑	↗	←	←
↙	↑	↗	↗	↗
↙	↑	↗	↗	↗
↙	↑	↗	↗	↗

$\pi_*$

# Example Grid World



- A bot is required to traverse a grid of 4×4 dimensions to reach its goal (1 or 16)
- There are 2 terminal states (1 and 16) and 14 non-terminal states (2 to 15)
- Each step is associated with a reward of  $-1$
- Consider a random policy: at every state, the probability of every action {up, down, left, right} is 0.25
- Initialize  $v_1$  for the random policy with all 0s

# Example Grid World: Policy Evaluation

- Turning Bellman equation into an update:

$$v_{k+1}(s) \doteq \mathbb{E}_{\pi}[R_{t+1} + \gamma v_k S_{t+1} | S_t = s]$$

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

$$\begin{aligned}
 v_1(6) &= \sum_{a \in \{u,d,l,r\}} \pi(a|6) \sum_{s',r} p(s',r|6,a) [r + \gamma v_0(s')] \\
 &= \sum_{a \in \{u,d,l,r\}} \underbrace{\pi(a|6)}_{= 0.25 \forall a} \sum_{s'} \underbrace{p(s'|6,a)}_{= -1} \underbrace{[r + \gamma v_0(s')]}_{= 0 \forall s'} \\
 &= 0.25 * \{-p(2|6,u) - p(10|6,d) - p(5|6,l) - p(7|6,r)\} \\
 &= 0.25 * \{-1 - 1 - 1 - 1\} \\
 &= -1 \\
 &\Rightarrow v_1(6) = -1
 \end{aligned}$$

- For non-terminal states,  $v_1(s) = -1$
- For terminal states,  $p(s'|s, a) = 0$  (and hence  $v_k(1) = v_k(6) = 0$ , for all  $k$ )

$v_1 =$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0



# Example Grid World: Policy Evaluation

- Step 2, with discount factor  $\gamma = 1$

$$\begin{aligned}
 v_2(6) &= \sum_{a \in \{u,d,l,r\}} \underbrace{\pi(a|6)}_{= 0.25 \forall a} \sum_{s'} p(s'|6,a) \underbrace{[r + \gamma v_1(s')]}_{= -1} = \begin{cases} -1, & s' \in S \\ 0, & s' \in S^+ \setminus S \end{cases} \\
 &= 0.25 * \{p(2|6,u)[-1 - \gamma] + p(10|6,d)[-1 - \gamma] \\
 &\quad + p(5|6,l)[-1 - \gamma] + p(7|6,r)[-1 - \gamma]\} \\
 &\stackrel{\gamma=1}{=} 0.25 * \{-2 - 2 - 2 - 2\} \\
 &= -2
 \end{aligned}$$

- For the other states (2, 5, 12, 15):

$$\begin{aligned}
 v_2(2) &= \sum_{a \in \{u,d,l,r\}} \underbrace{\pi(a|2)}_{= 0.25 \forall a} \sum_{s'} p(s'|2,a) \underbrace{[r + \gamma v_1(s')]}_{= -1} = \begin{cases} -1, & s' \in S \\ 0, & s' \in S^+ \setminus S \end{cases} \\
 &= 0.25 * \{p(2|2,u)[-1 - \gamma] + p(6|2,d)[-1 - \gamma] \\
 &\quad + p(1|2,l)[-1 - \gamma * 0] + p(3|2,r)[-1 - \gamma]\} \\
 &\stackrel{\gamma=1}{=} 0.25 * \{-2 - 2 - 1 - 2\} \\
 &= -1.75 \\
 &\Rightarrow v_2(2) = -1.75
 \end{aligned}$$

- For all red states,  $v_2(s) = -2$

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

$$v_2 =$$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

# Example Grid World: Policy Evaluation

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

Random policy

$k = 0$

0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0
0.0	0.0	0.0	0.0

$k = 1$

0.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	-1.0
-1.0	-1.0	-1.0	0.0

$k = 2$

0.0	-1.7	-2.0	-2.0
-1.7	-2.0	-2.0	-2.0
-2.0	-2.0	-2.0	-1.7
-2.0	-2.0	-1.7	0.0

$k = 3$

0.0	-2.4	-2.9	-3.0
-2.4	-2.9	-3.0	-2.9
-2.9	-3.0	-2.9	-2.4
-3.0	-2.9	-2.4	0.0

...

$k = 10$

0.0	-6.1	-8.4	-9.0
-6.1	-7.7	-8.4	-8.4
-8.4	-8.4	-7.7	-6.1
-9.0	-8.4	-6.1	0.0

...

$k = \infty$

0.0	-14.	-20.	-22.
-14.	-18.	-20.	-20.
-20.	-20.	-18.	-14.
-22.	-20.	-14.	0.0

$\leftarrow v_\pi$

	←	←	↙
↑	↖	↘	↓
↑	↖	↘	↓
↙	→	→	

Optimal policy

# Approximation

- Optimal policies are **computationally costly** to find – we can only approximate
  - In tasks with small, finite state sets: **tabular methods**
  - Otherwise: function approximation using a more compact parameterized function representation (*e.g.* using neural networks)

→ The online nature of RL allows us to *put more effort into learning to make decisions for frequently encountered states*

# Temporal-Difference Learning

- TD methods update estimates based on immediately observed reward and state
- Update rule:
$$V(S_t) \leftarrow V(S_t) + \alpha[R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$
- Because TD bases its update in part on an existing estimate, it is a *bootstrapping* method

## Tabular TD(0) for estimating $v_\pi$


```
Input: the policy  $\pi$  to be evaluated
Algorithm parameter: step size  $\alpha \in (0, 1]$ 
Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$ 

Loop for each episode:
  Initialize  $S$ 
  Loop for each step of episode:
     $A \leftarrow$  action given by  $\pi$  for  $S$ 
    Take action  $A$ , observe  $R, S'$ 
     $V(S) \leftarrow V(S) + \alpha[R + \gamma V(S') - V(S)]$ 
     $S \leftarrow S'$ 
  until  $S$  is terminal
```

# Temporal-Difference Learning

- TD vs Dynamic Programming methods
  - TD methods **do not require a model** of the environment's dynamics (rewards and next-state probability distributions)
- TD vs Monte Carlo methods
  - TD methods are naturally implemented in an **online, fully incremental fashion**, while MC methods must wait until the end of an episode
    - Useful if episodes are very long, or in continuing tasks (that have no episodes at all)
- Usually, TD methods converge faster than MC methods on stochastic tasks

# Sarsa: On-policy TD Control




- Update rule:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

- This rule uses every element of the quintuple of events  $(S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1})$

## Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$   
 Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$   
 Loop for each episode:  
   Initialize  $S$   
   Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)  
   Loop for each step of episode:  
     Take action  $A$ , observe  $R, S'$   
     Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)  
      $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$   
      $S \leftarrow S'; A \leftarrow A';$   
   until  $S$  is terminal



Sarsa

- Converges to optimal policy and action-value function if all state-action pairs are visited infinitely and policy converges to greedy (e.g. using  $\varepsilon$ -greedy with  $\varepsilon = 1/t$ )


# Q-learning: Off-policy TD Control

- Update rule:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[ R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

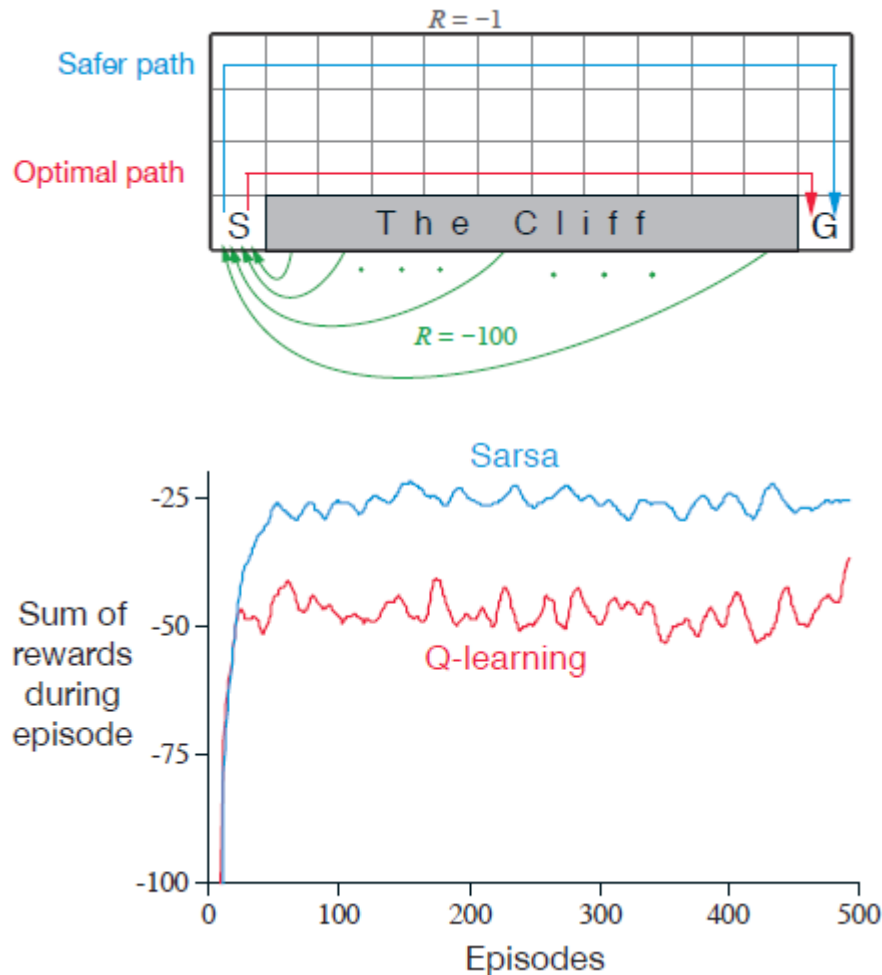
## Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$   
Initialize  $Q(s, a)$ , for all  $s \in S^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$   
Loop for each episode:  
  Initialize  $S$   
  Loop for each step of episode:  
    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)  
    Take action  $A$ , observe  $R, S'$   
     $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$   
     $S \leftarrow S'$   
  until  $S$  is terminal



- The learned action-value function  $Q$  directly approximates  $q_*$ , independently of the policy being followed

# Example: Cliff Walking




- Sarsa and Q-learning with  $\varepsilon$ -greedy action selection ( $\varepsilon = 0.1$ )
  - Q-learning learns values for the optimal policy
  - Sarsa takes action selection into account and learns the longer but safer path
  - Given exploration, Q-learning occasionally falls off the cliff, hence the lower online performance
- If  $\varepsilon$  is gradually reduced, both methods converge to the optimal policy



# Quiz: The Toothpick Game



- 10 toothpicks
- 2 players take 1, 2 or 3 in turn
- Avoid last toothpick
  - $r(0) = +1$        $r(1) = -1$        $r(n) = 0, n > 1$
- Using Q-learning with  $\alpha = 0.5$  and  $\gamma = 0.9$ , which values of  $Q(s, a)$  change in each of the following episodes? (Consider only self-actions, shown in **bold**.)
  - Actions: **2**-1-**2**-3-**1**-1 (States: **10**-8-**7**-5-**2**-1-0)
  - Actions: **2**-2-**1**-3-**1**-1 (States: **10**-8-**6**-5-**2**-1-0)
  - Actions: **1**-3-**1**-3-**1**-1 (States: **10**-9-**6**-5-**2**-1-0)
- After such updates, does the agent win when starting in state 10 and following a greedy policy, assuming the opponent always takes as much toothpicks as it can?




# RL Algorithms

Algorithm	Description	Policy	Action Space	State Space
Monte Carlo	Every visit to Monte Carlo	Either	Discrete	Discrete
Q-learning	State-action-reward-state	Off-policy	Discrete	Discrete
SARSA	State-action-reward-state-action	On-policy	Discrete	Discrete
Q-learning - Lambda	Q-learning with eligibility traces	Off-policy	Discrete	Discrete
SARSA - Lambda	SARSA with eligibility traces	On-policy	Discrete	Discrete
<a href="#">DQN</a> [Mnih et al., 2013]	Deep Q Network	Off-policy	Discrete	Continuous
<a href="#">DDPG</a> [Lillicrap et al., 2016]	Deep Deterministic Policy Gradient	Off-policy	Continuous	Continuous
<a href="#">A3C</a> [Mnih et al., 2016]	Asynchronous Advantage Actor-Critic	On-policy	Continuous	Continuous
<a href="#">NAF</a> [Gu et al., 2016]	Q-Learning with Normalized Advantage Functions	Off-policy	Continuous	Continuous
<a href="#">TRPO</a> [Schulman et al., 2015]	Trust Region Policy Optimization	On-policy	Continuous	Continuous
<a href="#">PPO</a> [Schulman et al., 2017]	Proximal Policy Optimization	On-policy	Continuous	Continuous
<a href="#">TD3</a> [Fujimoto et al., 2018]	Twin Delayed Deep Deterministic Policy Gradient	Off-policy	Continuous	Continuous
<a href="#">SAC</a> [Haarnoja et al., 2018]	Soft Actor-Critic	Off-policy	Continuous	Continuous

# RL in Games

- TD-Gammon [Tesauro, 1995]
  - Neural Network trained with self-play reinforcement learning





- Atari 2600 Games [DeepMind, 2013]
  - Learn control policies directly from high-dimensional sensory input using reinforcement learning
  - Input is raw pixels and output is a value function estimating future rewards




# RL in Games

- AlphaGo [Google DeepMind, 2016]
  - Convolutional Neural Networks trained with human expert data
  - Deep Reinforcement Learning with fictitious self-play




- Poker: Heads-Up Limit Texas Hold'em – NFSP [UCL, 2016]
  - Deep Reinforcement Learning with fictitious self-play
  - No prior knowledge



# OpenAI Gym

- [Gym](#) is a toolkit for developing and comparing reinforcement learning algorithms
- The [gym library](#) is a collection of test problems with a shared interface — **environments** — that you can use to work out your RL algorithms




(a) Toy Text



(b) Atari



(c) Controls



(d) MuJoCo



(e) Doom




(f) Minecraft

- [OpenAI Baselines](#) is a set of high-quality implementations of RL algorithms
  - See also [Stable Baselines](#)

# Unity ML-Agents

- With Unity Machine Learning Agents ([ML-Agents](#)), you teach intelligent agents through a combination of **deep reinforcement learning** and **imitation learning**



# Further Reading

- Sutton, R. S. and Barto, A. G. (2018). *Reinforcement Learning – An Introduction*, 2<sup>nd</sup> ed., The MIT Press: Chap. 1-3, 6
- Simple Tutorial Videos for Deep RL:
  - Introduction on Reinforcement Learning and Deep RL:  
<https://www.youtube.com/watch?v=JgvyzlkgxF0>
  - PPO – Proximal Policy Optimization (PPO):  
<https://www.youtube.com/watch?v=5P7I-xPq8u8>

# Conclusions

- RL enables to learn intelligent behavior in complex environments
- Large number of algorithms and approaches
- Amazing results in vintage Atari Games
- Stunning results of AlphaGo and AlphaZero
- Very promising results in Robotics
- Very fast evolution in the last few years